

Operating systems

Signals

Samuele Germiniani
samuele.germiniani@univr.it

University of Verona
Department of Engineering for Innovation Medicine

2023/2024



Table of Contents

1 Signals

- Fundamental concepts
- Signal types
- Signal handler
- Sending a signal
- Setting and blocking a signal



Signals



Signals

Fundamental concepts



Fundamental concepts (1/2)

A *signal* is a notification to a process that an event has occurred. They interrupt the normal flow of execution of a program; in most cases, it is not possible to predict exactly when a signal will arrive.

A signal is said to be *generated* by some event. Once generated, a signal is later *delivered* to a process. Between the time it is generated and the time it is delivered, a signal is said to be *pending*.

Normally, a pending signal is delivered to a process as soon as it is next scheduled to run, or *immediately* if the process is already running.



Fundamental concepts (2/2)

Upon delivery of a signal, a process carries out one of the following default actions, depending on the signal:

- The process is terminated (*killed*).
- The process is suspended (*stopped*).
- The process is *resumed* after previously being stopped.
- The signal is ignored. It is discarded by the kernel and has no effect on the process. (The process never even knows that it occurred.)
- The process executes a *signal handler*, namely a function written by the programmer that performs appropriate tasks in response to the delivery of a signal.



Signals

Signal types

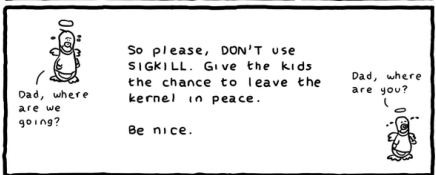






Signal Types and Default Actions (1/3)

Signals to terminate a process:

- **SIGTERM** is delivered to safely terminate a process. A well-designed application should have a handler for SIGTERM that causes the application to exit gracefully.
- **SIGINT** terminates a process (“interrupt process”). It is sent when the user type Contr-C character.
- **SIGQUIT** terminates a process and causes it to produce a core dump, which can then be used for debugging.
- **SIGKILL** terminates a process (always!). It can't be blocked, ignored, or caught by a handler.





Daniel Stori (turnoff.us)



Signal Types and Default Actions (2/3)

Signals to stop and resume a process:

- **SIGSTOP** stops a process (always!). It can't be blocked, ignored, or caught by a handler.
- **SIGCONT** resumes a previously stopped process.

Other important signals:

- **SIGPIPE** is generated when a process tries to write to a PIPE, a FIFO for which there is no corresponding reader process (see chapter PIPE/FIFO).
- **SIGALRM** is delivered to a process upon the expiration of a real-time timer set by a call to *alarm* (see next slides).
- **SIGUSR1** and **SIGUSR2** are available for programmer-defined purposes. The kernel never generates these signals for a process.

The complete list of available signals in Linux can be retrieved with the bash command “man 7 signal”.



Signal Types and Default Actions (3/3)

name	number	can be caught?	default action
SIGTERM	15	yes	terminates a process
SIGINT	2	yes	terminates a process
SIGQUIT	3	yes	dumps + terms a process
SIGKILL	9	no	terminates a process
SIGSTOP	17	no	stops a process
SIGCONT	19	yes	resumes a stopped process
SIGPIPE	13	yes	terminates a process
SIGALRM	14	yes	terminates a process
SIGUSR1	30	yes	terminates a process
SIGUSR2	31	yes	terminates a process

Column “number” reports the signal number for x86 and arm architecture.
A signal may have a different number in other architectures



Signals

Signal handler



Signal handler

A signal handler (also called a signal catcher) is a function that is called when a specified signal is delivered to a process. It has always the following general form:

```
void sigHandler(int sig) {  
    /* Code for the handler */  
}
```

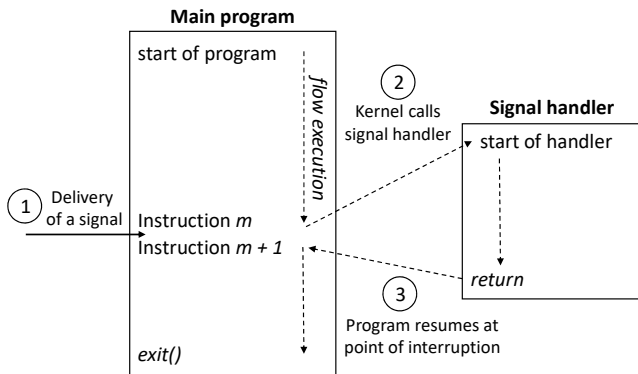
This function returns nothing (void) and takes one integer argument (*sig*). When the signal handler is invoked by the kernel, *sig* is set to the signal number delivered to the process.

Typically, *sig* is used to determine which signal caused the handler to be invoked when a same handler catches different types of signals.



Signal handler (execution)

Invocation of a signal handler may interrupt the main program flow at any time. The kernel calls the signal handler, and when the handler returns, execution of the program resumes at the point where the handler interrupted it.



Changing signal dispositions (1/2)

The *signal* system call changes the default signal-handler for a defined signal in a process.

```
#include <signal.h>

typedef void (*sighandler_t)(int);
// Returns previous signal disposition on success, or SIG_ERR on error
sighandler_t signal(int signum, sighandler_t handler);
```

signum identifies the signal whose disposition we wish to change in the process. *handler* can be one of the following:

- the address of a user-defined signal handler.
- the constant SIG_DFL, which resets the default disposition of the process for the signal *signum*.
- the constant SIG_IGN, which sets the process to ignore the delivery of the signal *signum*.



Changing signal dispositions (2/2)

```
void sigHandler(int sig) {
    printf("The signal %s was caught!\n",
        (sig == SIGINT)? "Ctrl-C" : "signal User-1");
}

int main (int argc, char *argv[]) {
    // setting sigHandler to be executed for SIGINT or SIGUSR1
    if (signal(SIGINT, sigHandler) == SIG_ERR ||
        signal(SIGUSR1, sigHandler) == SIG_ERR) {
        errExit("change signal handler failed");
    }
    // Do something else here. During this time, if SIGINT/SIGUSR1
    // is delivered, sigHandler will be used to handle the signal.
    // Reset the default process disposition for SIGINT and SIGUSR1
    if (signal(SIGINT, SIG_DFL) == SIG_ERR ||
        signal(SIGUSR1, SIG_DFL) == SIG_ERR) {
        errExit("reset signal handler failed");
    }
    return 0;
}
```



Signal handler (important notes)

What you should keep in mind when you use signal handlers:

- SIGKILL and SIGSTOP cannot be caught.
- A signal is an asynchronous event. We cannot predict when it arrives.
- When a signal handler is invoked, the signal that caused its invocation is automatically blocked. It is unblocked when the signal handler returns to the normal execution flow of the program.
- If a blocked signal is generated several times, when unblocked, it is delivered to the process only once!
- The execution of a signal handler can be interrupted by the delivery of an unblocked signal.
- The signal dispositions are inherited between process *parent* and process *child*.



Waiting for a signal (1/2)

Calling *pause* suspends execution of the process until the call is interrupted by a signal handler (or until an unhandled signal terminates the process).

```
#include <unistd.h>
// Always return -1 with errno set to EINTR
int pause();
```

The *sleep* function suspends execution of the calling process for the number of seconds specified in the *seconds* argument or until a signal is caught (thus interrupting the call).

```
#include <unistd.h>
unsigned int sleep(unsigned int seconds); // Returns 0 on normal completion, or
number of unslept seconds if prematurely terminated
```



Waiting for a signal (2/2)

Waiting the interrupt signal (Ctrl-C), which must occur within 30 seconds

```
void sigHandler(int sig) { printf("Well done!\n"); }

int main (int argc, char *argv[]) {
    if (signal(SIGINT, sigHandler) == SIG_ERR)
        errExit("change signal handler failed");

    int time = 30;
    printf("We can wait for %d seconds!\n", time);
    time = sleep(time); // the process is suspended for max. 30sec.
    printf("%s!\n", (time==0)? "out of time", "just in time");
}
```



Signals

Sending a signal



Sending a signal (kill) (1/4)

The system call *kill* let a process send a signal to another process.

```
#include <signal.h>

// Returns 0 on success, or -1 on error
int kill(pid_t pid, int sig);
```

The *pid* argument identifies one or more processes to which the signal specified by *sig* is to be sent.

- ($pid > 0$): the signal is sent to the process having PID equals to *pid*.
- ($pid = 0$): the signal is sent to every process in the same process group as the calling process, including the calling process itself.
- ($pid < 0$): the signal is sent to all of the processes in the process group whose ID equals the absolute value of *pid*.
- ($pid = -1$): the signal is sent to every process for which the calling process has permission to send a signal, except *init* and the process itself.



Sending a signal (kill) (2/4)

Sending a SIGKILL signal to a child process

```
int main (int argc, char *argv[]) {  
    pid_t child = fork();  
    switch(child) {  
        case -1:  
            errExit("fork");  
        case 0: /* Child process */  
            while(1); // <- child is stuck here!  
        default: /* Parent process */  
            sleep(10);           // wait 10 seconds  
            kill(child, SIGKILL); // kill the child process  
    }  
    return 0;  
}
```



Sending a signal (alarm) (3/4)

The *alarm* system call arranges for a SIGALRM signal to be delivered to the calling process after a fixed delay.

```
#include <signal.h>

// Always succeeds, returning number of seconds remaining on
// any previously set timer, or 0 if no timer previously was set
unsigned int alarm(unsigned int seconds);
```

- The *seconds* argument specifies the number of seconds in the future when the timer is to expire. At that time, a SIGALRM signal is delivered to the calling process.
- Setting a timer with *alarm* overrides any previously set timer.



Sending a signal (alarm) (4/4)

Setting a timer with the *alarm* system call.

```
void sigHandler(int sig) { printf("Out of time!\n"); _exit(0); }

int main (int argc, char *argv[]) {
    if (signal(SIGALRM, sigHandler) == SIG_ERR)
        errExit("change signal handler failed");

    int time = 30;
    printf("We have %d seconds to complete the job!\n", time);
    alarm(time); // setting a timer

    /* Do something else here. */

    time = alarm(0); // disabling timer
    printf("%s seconds before timer expirations!\n", time);
    return 0;
}
```



Signals

Setting and blocking a signal



Signal set (1/2)

The *sigset_t* data type represents a signal set. The functions *sigemptyset* and *sigfillset* must be used to initialize a signal set, before using it in any other way.

```
#include <signal.h>

typedef unsigned long sigset_t;

// Both return 0 on success, or -1 on error.
int sigemptyset(sigset_t *set);
int sigfillset(sigset_t *set);
```

sigemptyset initializes a signal set to contain no signal.
sigfillset initializes a set to contain all signals.



Signal set (2/2)

After initialization, individual signals can be added to a set using *sigaddset* and removed using *sigdelset*.

```
#include <signal.h>

// Both return 0 on success, or -1 on error
int sigaddset(sigset_t *set, int sig);
int sigdelset(sigset_t *set, int sig);
```

For both *sigaddset* and *sigdelset*, the *sig* argument is a signal number.

The *sigismember* function is used to test for membership of a set.

```
#include <signal.h>

// Returns 1 if sig is a member of set, otherwise 0
int sigismember(const sigset_t *set, int sig);
```



Blocking signal delivery (1/3)

For each process, the kernel maintains a signal mask, namely a set of signals whose delivery to the process is currently blocked. If a signal that is blocked is sent to a process, delivery of that signal is delayed until it is unblocked by being removed from the process signal mask.

The *sigprocmask* system call can be used at any time to explicitly add signals to, and remove signals from, the signal mask.

```
#include <signal.h>

// Returns 0 on success, or -1 on error
int sigprocmask(int how, const sigset_t *set, sigset_t *oldset);
```



Blocking signal delivery (2/3)

The *how* argument determines the changes that *sigprocmask* makes to the signal mask:

- **SIG_BLOCK** The set of blocked signals is the union of the current set and the *set* argument.
- **SIG_UNBLOCK** The signals in *set* argument are removed from the current set of blocked signals. It is permissible to attempt to unblock a signal which is not blocked.
- **SIG_SETMASK** The set of blocked signals is set to the argument *set*.

In each case, if the *oldset* argument is not NULL, it points to a *sigset_t* buffer that is used to return the previous signal mask.

If we want to retrieve the signal mask without changing it, then we can specify NULL for the *set* argument, in which case the *how* argument is ignored.



Blocking signal delivery (3/3)

Blocking all signals but SIGTERM.

```
int main (int argc, char *argv[]) {
    sigset_t mySet, prevSet;
    // initialize mySet to contain all signals
    sigfillset(&mySet);
    // remove SIGTERM from mySet
    sigdelset(&mySet, SIGTERM);
    // blocking all signals but SIGTERM
    sigprocmask(SIG_SETMASK, &mySet, &prevSet);

    // the process is not interrupted by signals except SIGTERM

    // reset the signal mask of the process
    sigprocmask(SIG_SETMASK, &prevSet, NULL);
    // the process is not interrupted by signals in prevSet
    return 0;
}
```

