

Operating systems

Semaphores (IPC)

Samuele Geminiani
samuele.germiniani@univr.it

University of Verona
Department of Engineering for Innovation Medicine

2023/2024



Table of Contents

1 Semaphores

- Creating and Opening
- Control Operations
- Other Operations



Semaphores



Semaphores

Creating and Opening



Creating/Opening a Semaphore Set

The `semget` system call creates a new **semaphore set** or obtains the identifier of an existing set.

```
#include <sys/sem.h>

// Returns semaphore set identifier on success, or -1 error
int semget(key_t key, int nsems, int semflg);
```

The key arguments are: an IPC key, `nsems` specifies the number of semaphores in that set, and must be greater than 0. `semflg` is a bit mask specifying the permissions (see `open(...)` system call, `mode` argument) to be placed on a new semaphore set or checked against an existing set.

In additions, the following flags can be ORed (`|`) in `semflg`:

- `IPC_CREAT`: If no semaphore set with the specified key exists, create a new set.
- `IPC_EXCL`: in conjunction with `IPC_CREAT`, it makes `semget` fail if a semaphore set exists with the specified key.



Creating/Opening a Semaphore Set

Example showing how to create a semaphore set having 10 semaphores

```
int semid;
key_t key = //... (generate a key in some way, i.e. with ftok)

// A) delegate the problem of finding a unique key to the kernel
semid = semget(IPC_PRIVATE, 10, S_IRUSR | S_IWUSR);

// B) create a semaphore set with identifier key, if it doesn't already exist
semid = semget(key, 10, IPC_CREAT | S_IRUSR | S_IWUSR);

//C) create a semaphore set with identifier key, but fail if it exists already
semid = semget(key, 10, IPC_CREAT | IPC_EXCL | S_IRUSR | S_IWUSR);
```



Semaphore Control Operations

The `semctl` system call performs a variety of control operations on a semaphore set or on an individual semaphore within a set.

```
#include <sys/sem.h>

// Returns nonnegative integer on success, or -1 error
int semctl(int semid, int semnum, int cmd, ... /* union semun arg */);
```

The `semid` argument is the identifier of the semaphore set on which the operation is to be performed.

Certain control operations (`cmd`) require a third/fourth argument. Before presenting the available control operations on a semaphore set, the union `semun` is introduced.



Semaphore Control Operations - union *semun*

The union `semun` must be **explicitly defined by the programmer** before calling the `semctl` system call.

```
#ifndef SEMUN_H
#define SEMUN_H
#include <sys/sem.h>
// definition of the union semun
union semun {
    int val;
    struct semid_ds * buf;
    unsigned short * array;
};
#endif
```



Semaphores

Control Operations



Semaphore Control Operations

Generic control operations

```
int semctl(semid, 0 /*ignored*/, cmd, arg);
```

- IPC_RMID: Immediately remove the semaphore set. Any processes blocked is awakened (error set to EIDRM). The arg argument is not required.
- IPC_STAT: Place a copy of the `semid_ds` data structure associated with this semaphore set in the buffer pointed to by `arg.buf`.
- ICP_SET: Update selected fields of the `semid_ds` data structure associated with this semaphore set using values in the buffer pointed to by `arg.buf`.



Semaphore Control Operations

Generic control operations

```
struct semid_ds {  
    struct ipc_perm sem_perm; /* Ownership and permissions */  
    time_t sem_otime;         /* Time of last semop() */  
    time_t sem_ctime;         /* Time of last change */  
    unsigned long sem_nsems; /* Number of semaphores in set */  
};
```

Only the subfields *uid*, *gid*, and *mode* of the substructure *sem_perm* can be updated via `IPC_SET`.



Semaphore Control Operations

Generic control operations (Example)

Example showing how to **change the permissions** of a semaphore set

```
ket_t key = //... (generate a key in some way, i.e. with ftok)
// get, or create, the semaphore set
int semid = semget(key, 10, IPC_CREAT | S_IRUSR | S_IWUSR);
// instantiate a semid_ds struct
struct semid_ds ds;
// instantiate a semun union (defined manually somewhere)
union semun arg;
arg.buf = &ds;
// get a copy of semid_ds structure belonging to the kernel
if (semctl(semid, 0 /*ignored*/, IPC_STAT, arg) == -1)
    errExit("semctl IPC_STAT failed");
// update permissions to guarantee read access to the group
arg.buf->sem_perms.mode |= S_IRGRP;
// update the semid_ds structure of the kernel
if (semctl(semid, 0 /*ignored*/, IPC_SET, arg) == -1)
    errExit("semctl IPC_SET failed");
```



Semaphore Control Operations

Generic control operations (Example)

Example showing how to **remove** semaphore set

```
if (semctl(semid, 0/*ignored*/, IPC_RMID, 0/*ignored*/) == -1)
    errExit("semctl failed");
else
    printf("semaphore set removed successfully\n");
```



Semaphore Control Operations

Retrieving and initializing semaphore values

```
int semctl(semid, semnum, cmd, arg);
```

- SETVAL: the value of the *semnum-th* semaphore in the set referred to by *semid* is initialized to the value specified in *arg.val*.
- GETVAL: as its function result, *semctl* returns the value of the *semnum-th* semaphore in the semaphore set specified by *semid*. The *arg* argument is not required.

```
int semctl(semid, 0 /*ignored*/, cmd, arg);
```

- SETALL: initialize all semaphores in the set referred to by *semid*, using the values supplied in the array pointed to by *arg.array*.
- GETALL: retrieve the values of all of the semaphores in the set referred to by *semid*, placing them in the array pointed to by *arg.array*.



Semaphore Control Operations

Retrieving and initializing semaphore values (Example)

Example showing how to **initialize a specific semaphore** in a semaphore set

```
key_t key = //... (generate a key in some way, i.e. with ftok)
// get, or create, the semaphore set
int semid = semget(key, 10, IPC_CREAT | S_IRUSR | S_IWUSR);
// set the semaphore value to 0
union semun arg;
arg.val = 0;
// initialize the 5-th semaphore to 0
if (semctl(semid, 5, SETVAL, arg) == -1)
    errExit("semctl SETVAL");
```

A semaphore set must be always initialized before using it!



Semaphore Control Operations

Retrieving and initializing semaphore values (Example)

Example showing how to **get the current state** of a specific semaphore in a semaphore set.

```
key_t key = //... (generate a key in some way, i.e. with ftok)
// get, or create, the semaphore set
int semid = semget(key, 10, IPC_CREAT | S_IRUSR | S_IWUSR);

// get the current state of the 5-th semaphore
int value = semctl(semid, 5, GETVAL, 0/*ignored*/);
if (value == -1)
    errExit("semctl GETVAL");
```

Once returned, the semaphore may already have changed state!



Semaphore Control Operations

Retrieving and initializing semaphore values (Example)

Example showing how to **initialize a semaphore** set having 10 semaphores

```
ket_t key = //... (generate a key in some way, i.e. with ftok)
// get, or create, the semaphore set
int semid = semget(key, 10, IPC_CREAT | S_IRUSR | S_IWUSR);
// set the first 5 semaphores to 1, and the remaining to 0
int values[] = {1,1,1,1,1,0,0,0,0,0};
union semun arg;
arg.array = values;
// initialize the semaphore set
if (semctl(semid, 0/*ignored*/, SETALL, arg) == -1)
    errExit("semctl SETALL");
```

A semaphore set must be always initialized before using it!



Semaphore Control Operations

Retrieving and initializing semaphore values (Example)

Example showing how to **get the current state** of a semaphore set having 10 semaphores

```
ket_t key = //... (generate a key in some way, i.e. with ftok)
// get, or create, the semaphore set
int semid = semget(key, 10, IPC_CREAT | S_IRUSR | S_IWUSR);
// declare an array big enough to store the semaphores' value
int values[10];
union semun arg;
arg.array = values;
// get the current state of a semaphore set
if (semctl(semid, 0/*ignored*/, GETALL, arg) == -1)
    errExit("semctl GETALL");
```

Once returned, a semaphore may already have changed state!



Semaphore Control Operations

Retrieving per-semaphore information

```
int semctl(semid, semnum, cmd, 0);
```

- GETPID: return the process ID of the last process to perform a semop on the *semnum-th* semaphore
- GETNCNT: return the number of processes currently waiting for the value of the *semnum-th* semaphore to increase
- GETZCNT: return the number of processes currently waiting for the value of the *semnum-th* semaphore to become 0;



Semaphore Control Operations

Retrieving per-semaphore information (Example)

Example showing how to **get information about a semaphore** of the semaphore set

```
key_t key = //... (generate a key in some way, i.e. with ftok)
// get, or create, the semaphore set
int semid = semget(key, 10, IPC_CREAT | S_IRUSR | S_IWUSR);
// ...
// get information about the first semaphore of the semaphore set
printf("Sem:%d getpid:%d getncnt:%d getzcnt:%d\n",
semid,
semctl(semid, 0, GETPID, NULL),
semctl(semid, 0, GETNCNT, NULL),
semctl(semid, 0, GETZCNT, NULL));
```

Once returned, the semaphore may already have changed state!



Semaphores

Other Operations



Semaphore Operations

The `semop` system call performs one or more operations (wait (P) and signal (V)) on semaphores.

```
#include <sys/sem.h>

// Returns 0 on success, or -1 on error
int semop(int semid, struct sembuf *sops, unsigned int nsops);
```

The `sops` argument is a pointer to an array that contains a sorted sequence of operations to be performed atomically, and `nsops` (> 0) gives the size of this array. The elements of the `sops` array are structures of the following form:

```
struct sembuf {
    unsigned short sem_num; /* Semaphore number */
    short sem_op;           /* Operation to be performed */
    short sem_flg;         /* Operation flags */
};
```



Semaphore Operations

The `sem_num` field identifies the semaphore within the set upon which the operation is to be performed. The `sem_op` field specifies the operation to be performed:

- `sem_op > 0`: value of `sem_op` is added to the value of the `sem_num`-th semaphore.
- `sem_op = 0`: the value of the `sem_num`-th semaphore is checked to see whether it currently equals 0. If it doesn't, the calling process is blocked until the semaphore is 0.
- `sem_op < 0`: decrease the value of the `sem_num`-th semaphore by the amount specified in `sem_op`. it blocks the calling process until the semaphore value has been increased to a level that permits the operation to be performed without resulting in a negative value.



Semaphore Operations

When a `semop(...)` call blocks, the process remains blocked until one of the following occurs:

- Another process modifies the value of the semaphore such that the requested operation can proceed.
- A signal interrupts the `semop(...)` call. In this case, the error `EINTR` results.
- Another process deletes the semaphore referred to by `semid`. In this case, `semop(...)` fails with the error `EIDRM`.

We can prevent `semop(...)` from blocking when performing an operation on a particular semaphore by specifying the `IPC_NOWAIT` flag in the corresponding `sem_flg` field. In this case, if `semop(...)` would have blocked, it instead fails with the error `EAGAIN`.



Semaphore Operations

Example showing how to initialize an array of sembuf operations

```
struct sembuf sops[3];

sops[0].sem_num = 0;
sops[0].sem_op = -1; // subtract 1 from semaphore 0
sops[0].sem_flg = 0;

sops[1].sem_num = 1;
sops[1].sem_op = 2; // add 2 to semaphore 1
sops[1].sem_flg = 0;

sops[2].sem_num = 2;
sops[2].sem_op = 0; // wait for semaphore 2 to equal 0
sops[2].sem_flg = IPC_NOWAIT; // but don't block if operation cannot be
    performed immediately
```



Semaphore Operations

Example showing how to perform operations on a semaphore set

```
struct sembuf sops[3];

// .. see the previous slide to initialize sembuf

if (semop(semid, sops, 3) == -1) {
    if (errno == EAGAIN) // Semaphore 2 would have blocked
        printf("Operation would have blocked\n");
    else
        errExit("semop"); // Some other error
}
```

